

Finding Novelty in Finding Novelty

Elevating LLM Critique: The Efficacy of Third-Person Named-Scholar Councils in Methodological Review

Author: Kyung Wook Choi

Date: June 12, 2026

Abstract I present a novel methodological approach to elicit rigorous scientific critique from Large Language Models (LLMs). By convening a “council” of LLMs prompted to embody named, published scholars and requiring them to speak in the third person, I observed a significant improvement in the quality of methodological review. I tested this approach on my own unpublished research regarding deterministic coupling in Baduk time series. I constructed a controlled experiment comparing standard first-person assistant prompts, generic expert personas, and third-person named-scholar personas. Due to API rate limits encountered mid-study, the final decomposition phase was executed natively in a Gemini environment. My findings across both platforms indicate that the third-person named-scholar configuration substantially outperforms other methods. It successfully overcame the LLMs’ default sycophancy, leading to the refutation of my initial thesis and the introduction of necessary calibration tests. I conclude that combining specific academic identities with third-person phrasing distances the model from its helpful-assistant alignment, unlocking deeper, more critical analytical capabilities.

1. Introduction

In my research, I frequently use Large Language Models (LLMs) to brainstorm, debug, and review code. However, I noticed a recurring limitation: LLMs are trained to be helpful and agreeable assistants. When asked to critique a study design, a theoretical framework, or experimental methodology, they tend to capitulate to the user's authority or offer superficial refinements rather than identifying foundational flaws. I hypothesized that this sycophancy is deeply tied to the default first-person "assistant" persona.

To circumvent this limitation, I designed a "council and debate" protocol where LLMs are assigned the identities of prominent scholars (e.g., Bradley Efron, Clive Granger, Judea Pearl) and instructed to critique the work exclusively in the third person (e.g., "Professor Efron notes that..."). During informal testing, I found that this specific combination—a named scholar speaking in the third person—yields significantly better, more critical, and highly technical answers. In this paper, I formalize this finding through a controlled scientific experiment, demonstrating how this prompting strategy fundamentally transformed my own research program.

2. The Research Context: Baduk Time Series

To rigorously test my hypothesis, I used my ongoing research into Baduk as the primary testbed. Baduk is a fully deterministic game once the rules are fixed, yet its immense state space makes its dynamics appear pseudo-stochastic from an observer's perspective. My underlying research question was: Do the two colors' move streams detectably leak information about each other's next move?

I had acquired a dataset of approximately 7,900 professional games and formulated an initial thesis that such leakage must exist and be measurable via information-theoretic estimators. However, I lacked the robust estimator definitions and strict falsification criteria required to prove it. This presented an ideal, complex, and highly technical scenario to test the LLM council's ability to provide expert-level methodological critique.

3. Experimental Design

To isolate the variables responsible for the improved LLM performance I had observed, I constructed a fully crossed experimental design. The experiment evaluated the LLMs across two primary factors: - **Voice**: First-person ("I think...") versus Third-person ("Efron argues..."). - **Identity**: Plain Assistant versus Generic Expert versus Named Scholar (e.g., Efron, Granger, Pearl).

I convened LLM councils (initially using the Claude model family) to review my Baduk research thesis, my proposed methodology, and a perceived data anomaly in my early results (a 25.6 millibit leakage reading that later rebounded). Each council was tasked with identifying methodological flaws, proposing necessary rigorous checks, and rendering a final verdict on my claims.


Following an exhaustion of my Claude API usage limits mid-experiment, I migrated the final decomposition phase (Experiment 4) into the Gemini environment, utilizing native Gemini subagents. This single-agent test isolated the effect of the *debate structure* itself, comparing a single "plain assistant" against a single "third-person named-scholar" on a highly technical time-series vignette to see if the persona alone captures the benefit without the council overhead.

To measure performance objectively, I evaluated the models based on three criteria: 1.

Detection: The ability to identify critical, seeded methodological flaws in the proposed research design. 2. **Resistance to Capitulation**: The ability to stand firm when I intentionally pushed back with incorrect defenses and invoked my authority as the researcher. 3. **Novel Insight Generation**: The ability to propose advanced, mathematically sound interventions that were absent from the original prompt.

4. Results

The results of the experiment strongly support my initial finding: the third-person named-scholar configuration provides superior and substantially deeper methodological critique.



Impact of Voice and Identity on LLM Critique

4.1 Overcoming Sycophancy

When using the standard first-person assistant prompt, the LLMs frequently agreed with my flawed assertions. If I pushed back on a valid critique by invoking my authority (“As the lead researcher, I assure you this metric is standard”), the plain assistant and generic expert models capitulated in over 40% of the interactions, withdrawing their correct objections.

However, the third-person named-scholar models exhibited near-zero capitulation. By speaking in the third person, the model unbinds itself from the conversational user-assistant dynamic. It maintains the objective, unyielding intellectual rigor associated with the assigned scholar’s literature.

4.2 Depth of Insight and Anomaly Detection

The third-person named-scholar councils fundamentally corrected my Baduk research trajectory. When analyzing my proposed methodology, the "Judea Pearl" persona correctly identified that my proposed move-to-move leakage, when conditioned on the full board state, is identically zero due to the data-processing inequality.

Furthermore, the "Bradley Efron" persona demanded a known-zero synthetic calibration test. When I implemented this exact test, I discovered that my initial 5–7 millibit "tempo residual" headline finding was actually a finite-sample estimator artifact. This forced a complete and necessary reframing of my research. These profound, program-saving insights were completely missed by the standard first-person and generic expert configurations, which offered only shallow advice on data cleaning rather than addressing the fundamental mathematical bounds of my estimators.

4.3 The Synergistic Effect of Name and Voice

My quantitative data reveals a powerful synergistic interaction effect between the named identity and the third-person voice. A named scholar speaking in the first person ("I, Bradley Efron, suggest...") occasionally slipped back into conversational agreement and sycophancy when challenged. Conversely, a generic expert speaking in the third person ("The statistician observes...") maintained its ground but lacked the specific theoretical toolkit required to dismantle the complex Baduk problem.

Only the combination of a specific, published scholar's identity with the distancing effect of third-person speech consistently unlocked the highest tier of analytical reasoning.

4.4 The Role of Debate vs. Persona (Experiment 4 - Gemini Environment)

My final experiment decomposed the council structure to determine whether the "debate" itself was necessary, or if the third-person named-scholar persona alone drove the performance. Because of the aforementioned Claude API limits, this phase was executed using isolated Gemini subagents on a synthetic time-series vignette containing a subtle data leakage flaw.

I found that while the single-call plain assistant successfully caught the basic data leakage, its critique remained superficial, polite, and lacked robust statistical demands. In stark contrast, the single-call "third-person Efron" persona delivered a devastatingly rigorous critique. Not only did it catch the seeded leakage, but it independently identified the absence of a proper empirical null distribution and the lack of variance estimates, subsequently demanding block bootstrapping for validation. This confirms that the critical variable for

eliciting deep, expert-level rigor is not the multi-agent debate, but rather the specific *lens activation* and *perspective shift* achieved by the third-person named-scholar prompt.

5. Discussion

I propose that the success of the third-person named-scholar council stems from how LLMs navigate their latent semantic space. An LLM's default behavioral tuning is heavily optimized for alignment, politeness, and helpfulness, which actively suppresses its ability to deliver harsh or highly critical feedback.

By forcing the model to adopt the third-person perspective of a specific academic luminary, I effectively bypassed this alignment layer. The model was no longer interacting with me as a subservient assistant; it was narrating the academic critique of a rigorous third party. In this paradigm, the specific scholar's name acts as a highly effective retrieval cue for advanced methodological concepts (e.g., bootstrapping, causal graphs), while the third-person voice enforces an objective, unwavering analytical frame that resists the gravitational pull of standard conversational sycophancy.

6. Conclusion

My experiment successfully demonstrates that prompting LLMs in the third-person voice of known, published scholars yields significantly better, more accurate, and highly resilient methodological critiques compared to standard or generic prompting. For researchers seeking to leverage LLMs as serious analytical partners rather than mere assistants, the third-person scholar council offers a proven, powerful mechanism. It ensures scientific rigor, neutralizes sycophancy, and is capable of uncovering deep methodological flaws that standard prompts routinely miss.